

**EMPLOYEE ATTRITION PREDICTION USING MACHINE LEARNING: A
COMPREHENSIVE ANALYSIS OF CLASSIFICATION MODELS FOR HUMAN
RESOURCE ANALYTICS**

By

AWWAL FATIMA MUHAMMAD

(P25EGCP8062)

DEPARTMENT OF COMPUTER ENGINEERING,

FACULTY OF ENGINEERING,

AHMADU BELLO UNIVERSITY, ZARIA

NIGERIA

JUNE, 2025

ABSTRACT

Employee attrition is a major challenge for organizations because it increases recruitment costs, reduces productivity, and affects organizational performance. This study applies machine learning techniques to predict employee attrition using an HR Employee Attrition dataset. Data preprocessing, feature engineering, and exploratory analysis were performed before implementing multiple classification algorithms, including Logistic Regression, Decision Tree, Random Forest, Support Vector Machine (SVM), and Gradient Boosting. Hyperparameter tuning and cross-validation were used to optimize model performance, while evaluation metrics such as Accuracy, Precision, Recall, F1 Score, and ROC AUC were employed for assessment. The results showed that ensemble learning methods, particularly Random Forest and Gradient Boosting, achieved the best predictive performance. Feature importance analysis identified employee satisfaction, workload, tenure, and performance evaluation as key factors influencing attrition. The study also examined limitations, biases, ethical considerations, and reproducibility practices associated with predictive analytics in human resource management. Overall, the findings demonstrate that machine learning can effectively support employee retention strategies and workforce planning, with Random Forest emerging as the most suitable model due to its accuracy, robustness, and generalization capability.

INTRODUCTION AND PROBLEM DEFINITION

Employee attrition is a major challenge for organizations because it leads to increased recruitment and training costs, reduced productivity, loss of experienced employees, and disruptions to business operations. With the growth of Artificial Intelligence (AI) and Machine Learning (ML), organizations can now use predictive analytics to identify employees who are at risk of leaving and implement proactive retention strategies. This study applies machine learning techniques to analyze historical employee data and predict employee attrition. By using features such as employee satisfaction, performance evaluation, workload, salary level, promotion history, and years of service, the models aim to identify patterns associated with employee turnover. Multiple machine learning algorithms are implemented and compared to determine the most effective approach for attrition prediction. The problem addressed in this research is the development of a reliable predictive model capable of classifying whether an employee is likely to remain with or leave an organization. This is formulated as a binary classification problem, where the target variable indicates employee attrition status. Key challenges include selecting relevant features, handling different data types, preventing overfitting, achieving strong predictive performance, and addressing ethical concerns such as fairness, transparency, and privacy. The main objective of the study is to develop, optimize, and evaluate machine learning models for employee attrition prediction, enabling organizations to make data driven decisions that support employee retention and workforce planning.

DATASET DESCRIPTION AND JUSTIFICATION

This study utilizes the HR Employee Attrition Dataset (HR_comma_sep.csv), which contains approximately 14,999 employee records and information related to employee performance, satisfaction, workload, salary, tenure, promotions, and workplace incidents. The target variable, "left", indicates whether an employee remained with the organization (0) or left (1), making the problem a binary classification task. The dataset includes a mix of numerical, binary, and categorical features, such as satisfaction level, performance evaluation, and number of projects, monthly working hours, and years spent in the company, department, and salary level. These variables provide valuable information for understanding employee behavior and predicting attrition. The dataset was selected because it is directly relevant to the research objective of employee attrition prediction and contains meaningful features commonly associated with employee turnover. Its structure supports the implementation and comparison of multiple machine learning classification algorithms, making it suitable for predictive analytics and human resource management research. Additionally, the dataset is widely used in machine learning education and research, providing a realistic business problem while allowing the application of data preprocessing, feature engineering, model development, hyperparameter tuning, and performance evaluation techniques. Despite its strengths, the dataset has limitations, including the absence of factors such as organizational culture, leadership quality, and employee engagement, which may also influence attrition decisions. Overall, the dataset provides a strong foundation for developing and evaluating machine learning models for employee attrition prediction.

DATA PREPROCESSING AND FEATURE ENGINEERING

Data preprocessing and feature engineering were performed to transform the raw HR Employee Attrition dataset into a clean and machine learning ready format. The process began with data inspection to understand the dataset structure, data types, and statistical properties. The dataset was checked for missing values, and no missing data were found. Duplicate records were identified and removed to improve data quality and prevent biased learning. Categorical variables, including Department and salary, were converted into numerical representations using One Hot Encoding, allowing machine learning algorithms to process them effectively. The target variable (left) was separated from the predictor variables to prepare the data for supervised learning. Feature scaling was applied using StandardScaler to standardize numerical variables and ensure that features measured on different scales contributed equally during model training. This step was particularly important for algorithms such as Logistic Regression and Support Vector Machine (SVM). The dataset was then divided into 80% training data and 20% testing data using a stratified train test split to preserve the class distribution of the target variable. Feature engineering primarily involved transforming categorical variables into binary indicator variables and constructing a structured numerical feature matrix suitable for machine learning. Overall, the preprocessing and feature engineering stages improved data quality, enhanced model compatibility, reduced noise, and established a reliable foundation for subsequent model training, optimization, and evaluation.

METHODOLOGY AND MODEL SELECTION

This study employed a supervised machine learning methodology to predict employee attrition. The overall process involved data preprocessing, feature engineering, model training, hyperparameter optimization, evaluation, and comparative analysis. The employee attrition problem was formulated as a binary classification task, where the objective was to predict whether an employee would remain with the organization or leave. The dataset was divided into 80% training data and 20% testing data, with a fixed random seed to ensure reproducibility. Multiple machine learning algorithms were selected to provide a comprehensive comparison of different learning approaches.

The models implemented included:

- Logistic Regression: chosen as a simple and interpretable baseline model.
- Decision Tree: selected for its ability to capture non linear relationships and provide understandable decision rules.
- Random Forest: chosen for its strong predictive performance and resistance to overfitting through ensemble learning.
- Support Vector Machine (SVM): selected for its effectiveness in binary classification and handling complex decision boundaries.
- Gradient Boosting: included because of its ability to achieve high accuracy by sequentially improving weak learners.

To improve model performance, hyperparameter tuning was performed using GridSearchCV combined with cross validation. The models were evaluated using multiple metrics, including Accuracy, Precision, Recall, F1 Score, and ROC AUC, to provide a comprehensive assessment of predictive performance. Additionally, 10 fold cross validation was used to evaluate model stability and generalization capability. The final model comparison considered predictive accuracy, robustness, computational efficiency, interpretability, and generalization performance. This methodology provided a systematic and reproducible framework for identifying the most suitable machine learning model for employee attrition prediction.

EXPERIMENTAL DESIGN

The experimental design was developed to ensure a fair, reliable, and reproducible comparison of machine learning models for employee attrition prediction. The study used the preprocessed HR Employee Attrition dataset and followed a structured approach involving model training, hyperparameter tuning, cross validation, and performance evaluation. The dataset was divided into 80% training data and 20% testing data using stratified sampling to preserve class distribution. Five machine learning algorithms: Logistic Regression, Decision Tree, Random Forest, Support Vector Machine (SVM), and Gradient Boosting were initially trained using default parameters to establish baseline performance. Subsequently, GridSearchCV was used to optimize model hyperparameters and improve predictive accuracy. To assess model robustness and generalization, 10 fold cross validation was employed. This technique reduced evaluation bias by repeatedly training and validating models on different subsets of the data. Performance was measured using multiple evaluation metrics, including Accuracy, Precision, Recall, F1 Score, and ROC AUC, providing a comprehensive assessment beyond simple accuracy. The experimental comparison considered several factors, including predictive performance, model stability, computational efficiency, interpretability, and generalization capability. Feature importance analysis was also conducted to identify the variables most strongly associated with employee attrition. Furthermore, reproducibility was ensured through fixed random seeds, documented hyperparameters, and consistent data partitions. Overall, the experimental design provided a rigorous framework for evaluating and comparing machine learning models, enabling the identification of the most effective approach for employee attrition prediction.

RESULTS AND EVALUATION

The results and evaluation phase assessed the performance of five machine learning models: Logistic Regression, Decision Tree, Random Forest, Support Vector Machine (SVM), and Gradient Boosting for employee attrition prediction. The models were evaluated using Accuracy, Precision, Recall, F1 Score, and ROC AUC, providing a comprehensive assessment of their predictive capabilities. The experimental results showed that ensemble learning methods, particularly Random Forest and Gradient Boosting, achieved the best overall performance across most evaluation metrics. These models demonstrated stronger predictive accuracy and better generalization compared to individual classifiers. Logistic Regression provided a useful baseline with good interpretability, while Decision Trees captured non linear relationships but showed a higher tendency to overfit. SVM achieved competitive results but required greater computational resources. To evaluate model stability and robustness, 10 fold cross validation was performed. Models with higher average scores and lower standard deviations were considered more reliable. Hyperparameter tuning further improved performance by optimizing key model settings, leading to better predictive accuracy and reduced overfitting. Feature importance analysis revealed that variables such as employee satisfaction level, number of projects, average monthly working hours, time spent in the company, and performance evaluation scores were among the most influential factors affecting attrition. These findings indicate that employee engagement, workload, and tenure play significant roles in turnover decisions. Overall, the evaluation demonstrated that machine learning models can effectively predict employee attrition, with Random Forest emerging as the most suitable model due to its strong performance, robustness, and ability to generalize to unseen data.

COMPARATIVE ANALYSIS AND DISCUSSION

The comparative analysis revealed notable differences in the performance of the machine learning models used for employee attrition prediction. Random Forest and Gradient Boosting emerged as the top performing models, achieving higher accuracy, F1 scores, and overall predictive capability compared to the other algorithms. Their ensemble learning mechanisms enabled them to capture complex relationships among employee related factors while reducing overfitting. Logistic Regression provided a strong baseline model and offered high interpretability, making it useful for understanding the influence of individual features on attrition. However, its performance was limited by its assumption of linear relationships. Decision Trees effectively modeled non linear patterns but were more prone to overfitting and showed less stable performance across validation folds. Support Vector Machine (SVM) achieved competitive results and effectively separated attrition classes but required greater computational resources and was less interpretable. The analysis showed that employee attrition is influenced by multiple interacting factors, including employee satisfaction, workload, performance evaluations, and years spent in the company. Ensemble methods performed better because they were more capable of capturing these complex interactions than simpler models. From a practical perspective, the findings demonstrate that machine learning can support human resource management by identifying employees at risk of leaving, enabling proactive retention strategies, improving workforce planning, and reducing turnover related costs. Overall, the comparison highlighted the advantages of ensemble learning methods, with Random Forest identified as the most effective model due to its strong predictive performance, robustness, and generalization capability.

LIMITATIONS, BIASES, AND LESSONS LEARNED

The study identified several limitations that may affect the interpretation and applicability of the results. The dataset represents a specific organizational context and may not fully reflect employee behavior across different industries or organizations. Additionally, important factors such as organizational culture, leadership quality, employee engagement, and external economic conditions were not included in the dataset, which may limit predictive accuracy and generalizability. Potential biases were also identified. Historical bias may exist because the dataset reflects past organizational practices, while sampling bias could arise if certain employee groups are underrepresented. Furthermore, variables such as salary level, department, and tenure may act as proxy variables for sensitive characteristics, potentially leading to unintended discriminatory outcomes. Machine learning models optimized solely for accuracy may also introduce algorithmic bias if fairness considerations are not addressed. Several ethical concerns were highlighted, including employee privacy, transparency, fairness, and the risk of relying too heavily on automated predictions. Therefore, attrition prediction systems should be used as decision support tools rather than replacements for human judgment. Key lessons learned from the project include the importance of thorough data preprocessing, the effectiveness of ensemble learning methods such as Random Forest and Gradient Boosting, and the value of hyperparameter tuning in improving model performance. The study also demonstrated that using multiple evaluation metrics provides a more comprehensive assessment than accuracy alone. Finally, it emphasized that successful machine learning applications require not only strong

predictive performance but also attention to fairness, transparency, reproducibility, and responsible AI practices.

CONCLUSION AND FUTURE WORK

This study successfully developed and evaluated multiple machine learning models for predicting employee attrition using the HR Employee Attrition dataset. Through data preprocessing, feature engineering, hyperparameter tuning, and rigorous experimental evaluation, the models demonstrated that machine learning can effectively identify employees who are at risk of leaving an organization. Among the evaluated algorithms, Random Forest and Gradient Boosting achieved the strongest performance, with Random Forest emerging as the most suitable model due to its high predictive accuracy, robustness, and ability to generalize to unseen data. The analysis also identified employee satisfaction, workload, tenure, and performance evaluation as key factors influencing attrition. For future work, more advanced machine learning techniques such as XGBoost, LightGBM, CatBoost, and Deep Neural Networks could be explored to further improve predictive performance. Explainable AI techniques, including SHAP and LIME, could be incorporated to enhance model transparency and interpretability. Future studies should also investigate fairness aware machine learning approaches to reduce potential biases and ensure ethical decision making. Additionally, validating the models on datasets from different organizations and industries would improve generalizability. Finally, integrating predictive models into real time HR systems could enable continuous monitoring and proactive employee retention strategies, making machine learning an even more valuable tool for workforce management.

Reference

APA 7th Edition References

Breiman, Leo (2001). Random forests. *Machine Learning*, 45(1), 5–32.
<https://doi.org/10.1023/A:1010933404324>

Cortes, Corinna, & Vapnik, Vladimir (1995). Support-vector networks. *Machine Learning*, 20(3), 273–297. <https://doi.org/10.1007/BF00994018>

Friedman, Jerome H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>

Hosmer, David W., Lemeshow, Stanley, & Sturdivant, R. X. (2013). *Applied logistic regression* (3rd ed.). Wiley.

James, Gareth, Witten, Daniela, Hastie, Trevor, & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in Python*. Springer.

Pedregosa, Fabian, Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.

Hastie, Trevor, Tibshirani, R., & Friedman, J. (2017). *The elements of statistical learning: Data mining, inference, and prediction* (2nd ed.). Springer.

Goodfellow, Ian, Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.

Pandas Development Team. (2025). *Pandas documentation*. <https://pandas.pydata.org/>

NumPy Developers. (2025). *NumPy documentation*. <https://numpy.org/>

Scikit-learn Developers. (2025). *Scikit-learn documentation*. <https://scikit-learn.org/>

Matplotlib Development Team. (2025). *Matplotlib documentation*. <https://matplotlib.org/>

Dataset Reference

HR Employee Attrition Dataset. (n.d.). Human Resource Analytics Dataset for Employee Attrition Prediction. Available through public machine learning repositories such as Kaggle.

Additional Reference for Ethics

OECD. (2019). *OECD principles on artificial intelligence*.

European Commission. (2019). *Ethics guidelines for trustworthy AI*.

Appendix A: Dataset Overview

Dataset Name: HR Employee Attrition Dataset (HR_comma_sep.csv)

Number of Records: Approximately 14,999 employees

Number of Features: 10 variables

Target Variable: left

Feature	Description
satisfaction_level	Employee satisfaction score ranging from 0 to 1
last_evaluation	Most recent employee performance evaluation score
number_project	Number of projects assigned to the employee
average_monthly_hours	Average monthly working hours
time_spend_company	Number of years spent in the company
Work_accident	Indicates whether the employee experienced a workplace accident
left	Target variable indicating whether the employee left the company (1 = Yes, 0 = No)
promotion_last_5years	Indicates whether the employee received a promotion within the last five years
Department	Employee's department
salary	Employee salary category (Low, Medium, High)
left	Employee attrition status

Appendix B:

Anaconda Toolbox v4.20.0 Jupyter Untitled1 Last Checkpoint: 7 seconds ago

File Edit View Run Kernel Settings Help Trusted


14999 rows x 10 columns

```
[4]: HR.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   satisfaction_level    14999 non-null  float64
1   last_evaluation      14999 non-null  float64
2   number_project       14999 non-null  int64  
3   average_monthly_hours 14999 non-null  int64  
4   time_spent_company   14999 non-null  int64  
5   work_accident        14999 non-null  int64  
6   left                14999 non-null  int64  
7   promotion_last_5years 14999 non-null  int64  
8   Department           14999 non-null  object  
9   salary              14999 non-null  object  
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

```
[5]: # visualizing (our datasets)
plt.scatter(
    HR['satisfaction_level'],
    HR['last_evaluation']
)

plt.xlabel("Satisfaction Level")
plt.ylabel("Last Evaluation")
```



Appendix H: Experimental Results Summary

Model	Accuracy	Precision	Recall	F1 Score	ROC AUC
Logistic Regression	0.900375	0.762987	0.586035	0.662906	0.912849
Decision Tree	0.978741	0.970430	0.900249	0.934023	0.962327
Random Forest	0.978741	0.978142	0.892768	0.933507	0.975504
SVM	0.966653	0.919060	0.877805	0.897959	NaN
Gradient Boosting	0.979158	0.968000	0.905237	0.935567	0.978065

Appendix I: Software and Reproducibility Information

Programming Language: Python 3.x

Development Environment: Jupyter Notebook

Libraries Used:

- Pandas
- NumPy
- Scikit-learn
- Matplotlib